

NEURAL NETWORKS

Tan Kwan Chong

Chief Data Scientist, Booz Allen Hamilton

PROGRESSING IN YOUR DATA SCIENCE CAREER

LEARNING OBJECTIVES

- Understand various types of neural networks
- Applications of neural networks
- Apply a neural network model for regression
- Apply a neural network model for classification

OPENING

ARTIFICIAL NEURAL NETWORKS

OPENING

- Neural networks were first studied in the 1940s (!) as a model of biological neural networks
- Many advances since then have improved the ability to train and apply neural networks
- Good for both classification and regression but difficult to interpret model behaviors
- Deep learning in the past few years has been highly successful for otherwise difficult problems

OPENING

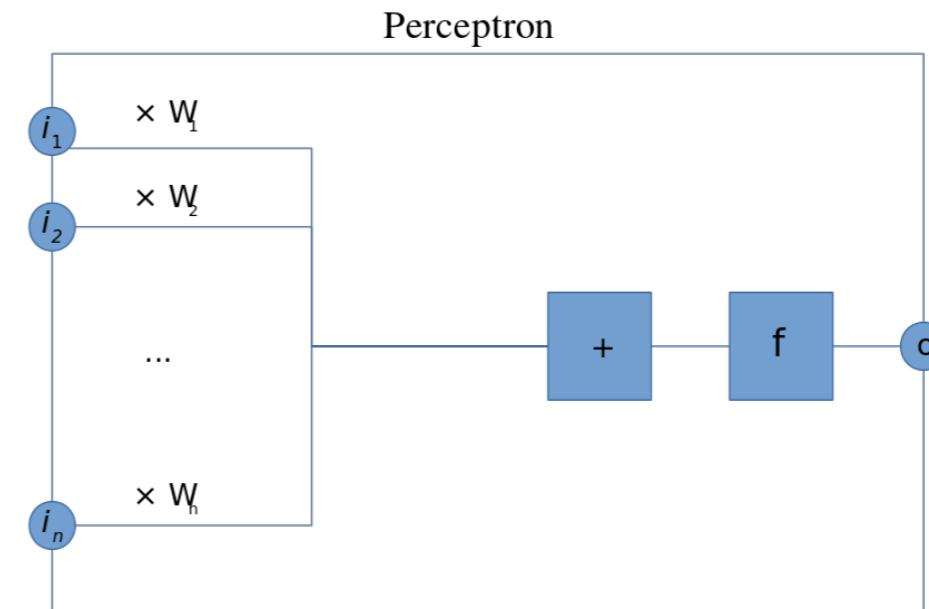
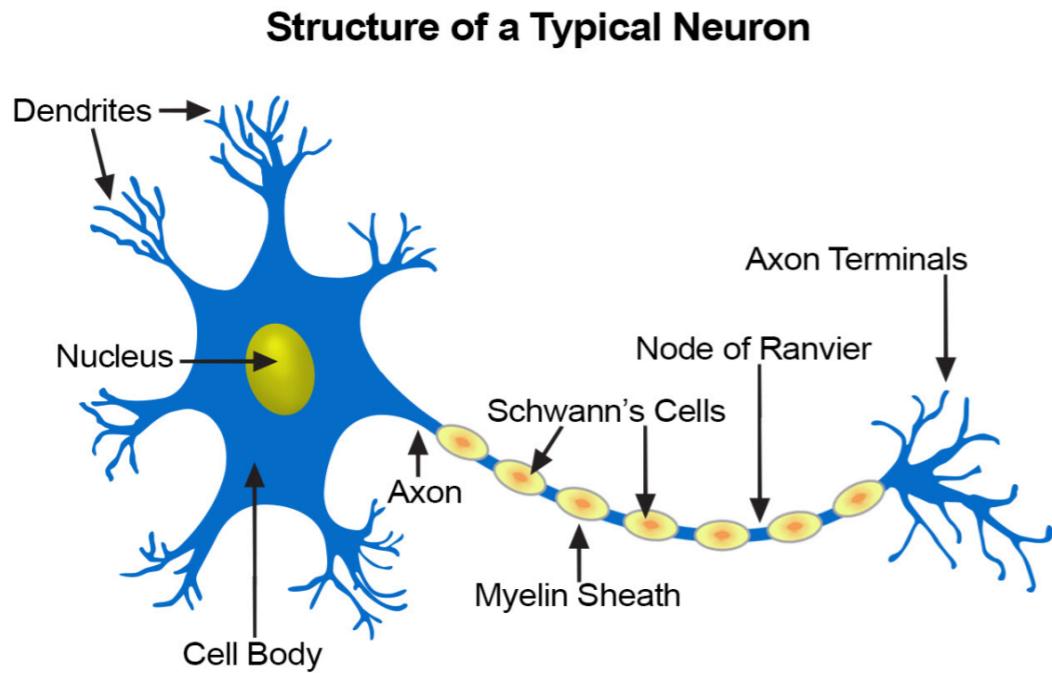
- Today we will focus on types of neural networks and their applications, and skip some of the more technical details
- Specifically we'll skip training neural networks -- there are many methods in various situations and the details can be tedious (but not particularly difficult)
- Methods include backpropagation, gradient descent, and Hessian-free learning

INTRODUCTION

PERCEPTRON

PERCEPTRON

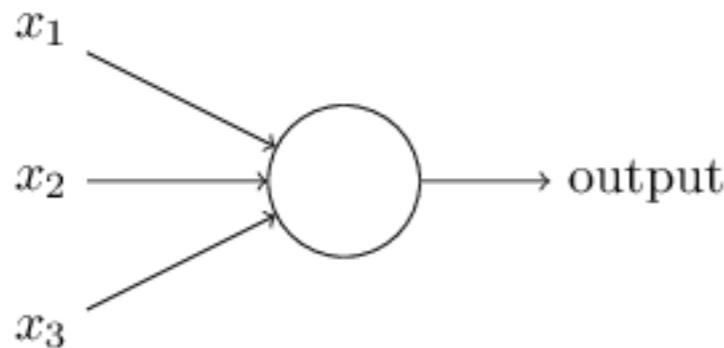
- Perceptrons are the simplest example of a neural network
- The idea is to emulate a single neuron



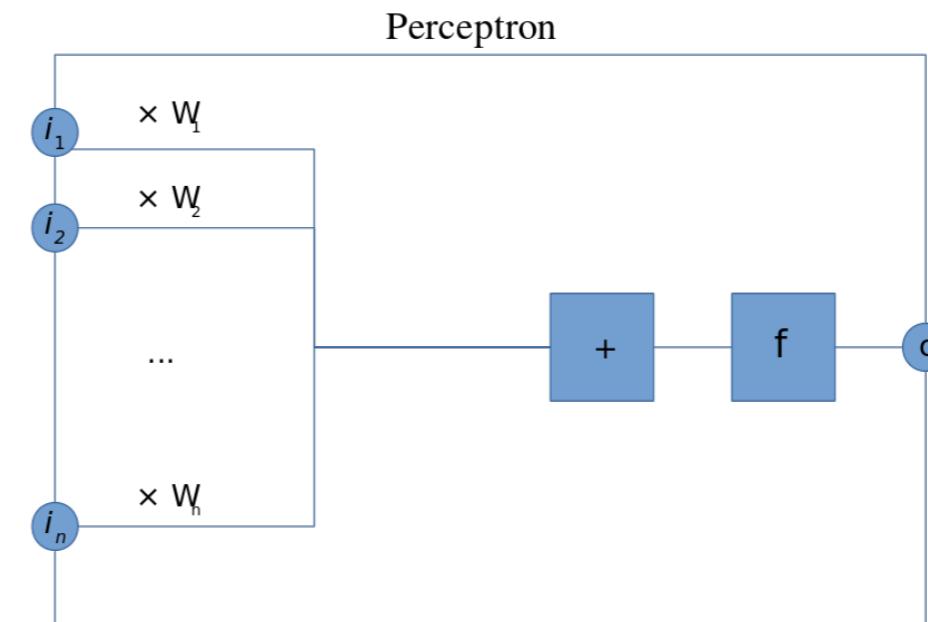
$$o = f\left(\sum_{k=1}^n i_k \cdot W_k\right)$$

PERCEPTRON

- Given n inputs and an activation or link function f
- A perceptron takes several binary inputs and produces a single binary output



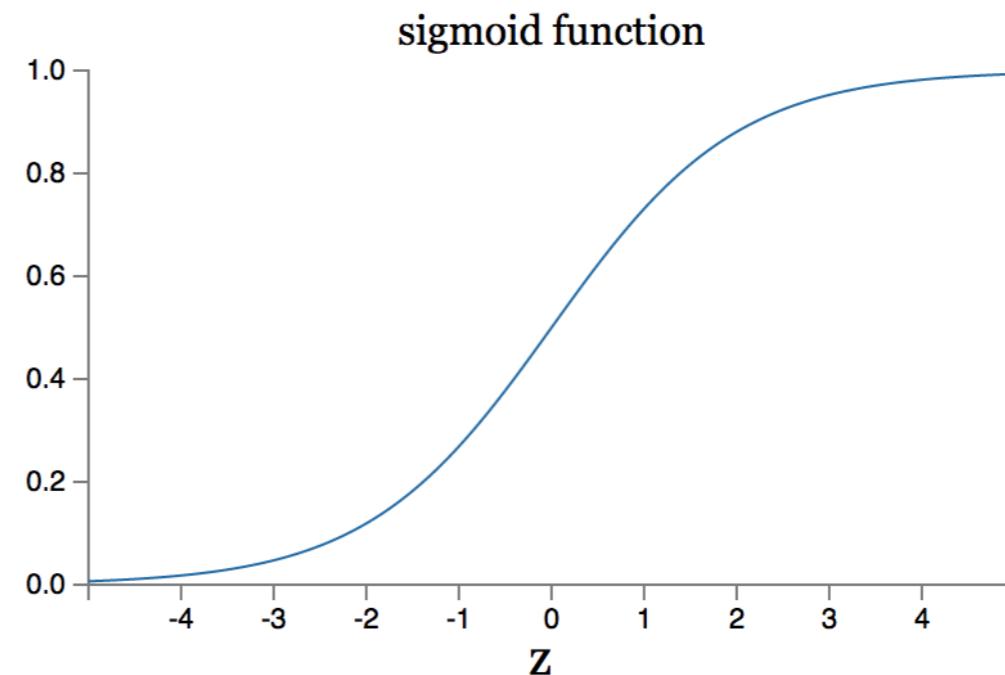
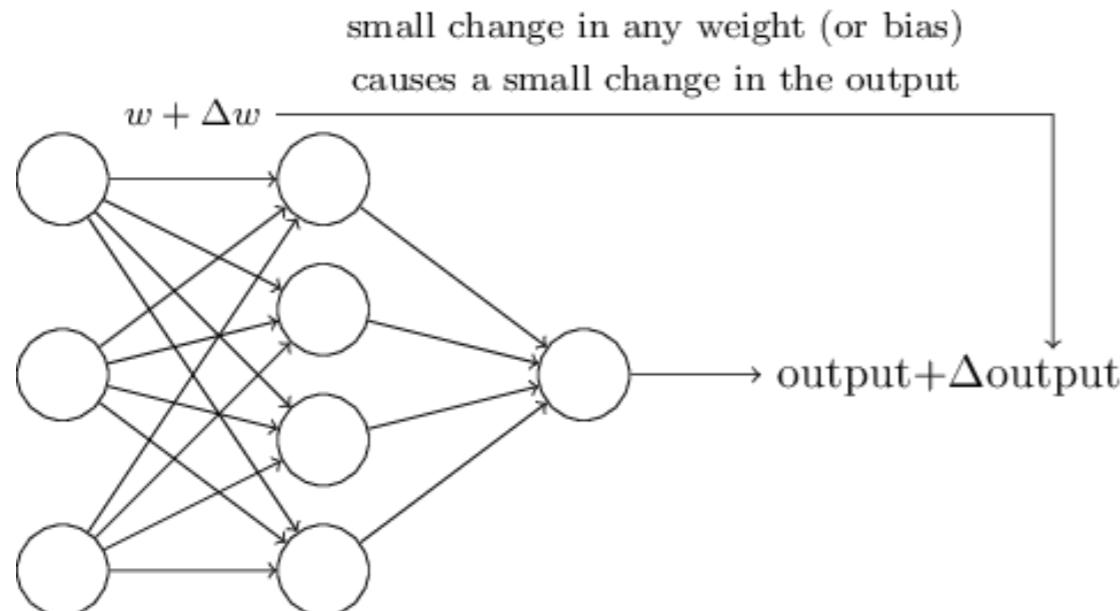
$$\text{output} = \begin{cases} 0 & \text{if } \sum_j w_j x_j \leq \text{threshold} \\ 1 & \text{if } \sum_j w_j x_j > \text{threshold} \end{cases}$$



$$o = f\left(\sum_{k=1}^n i_k \cdot W_k\right)$$

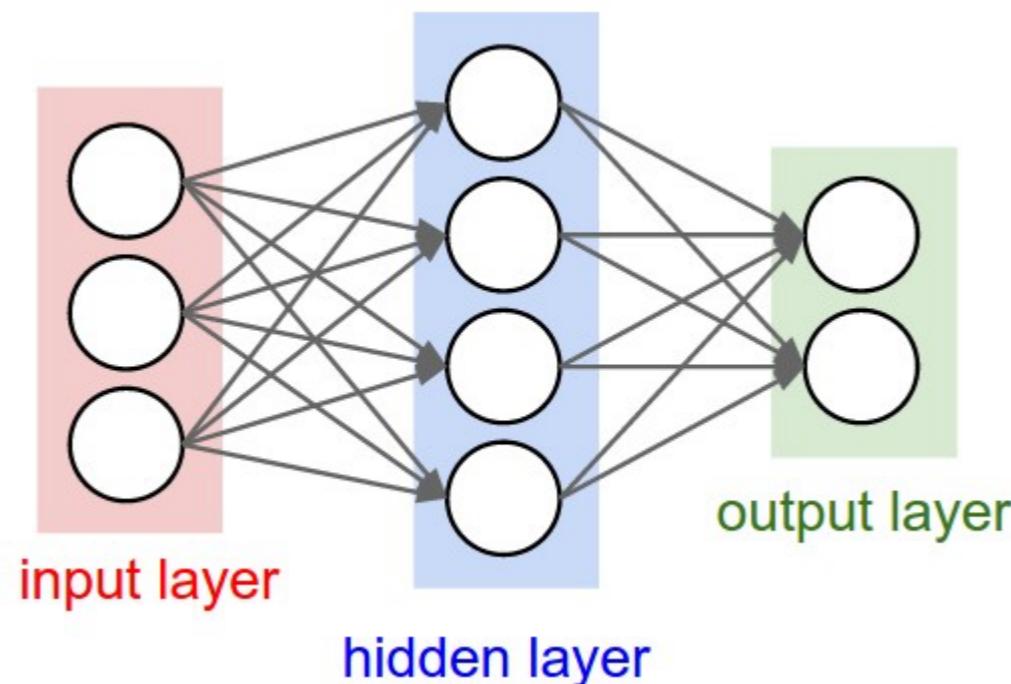
PERCEPTRON

- Common activation functions are linear, logistic, tanh, and softmax
- We'll see shortly that some are better for classification, some for regression
- Perceptrons can be combined into multilayer perceptrons or feed-forward network



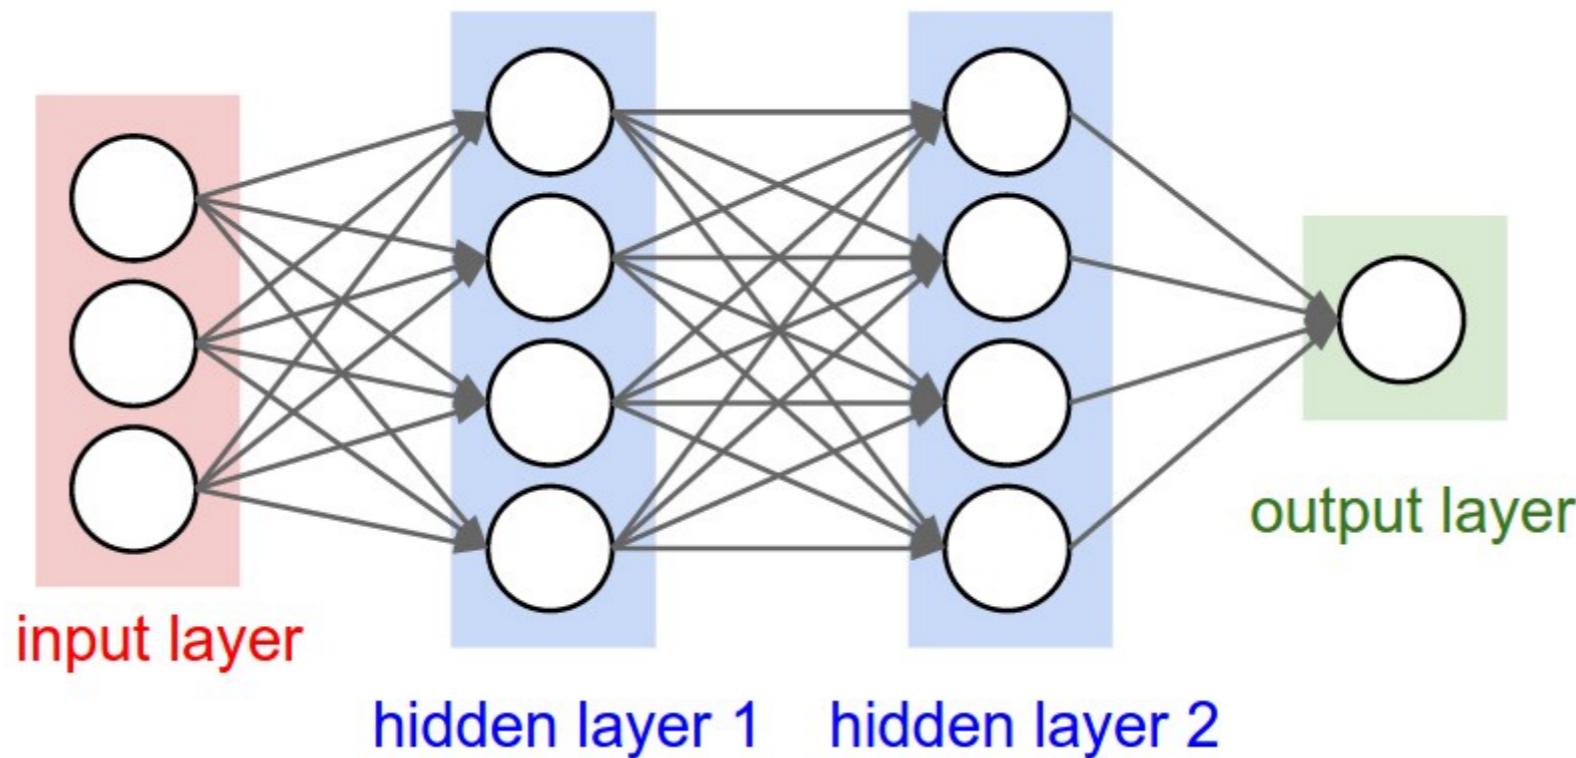
FEED FORWARD NN

- [Source](#)



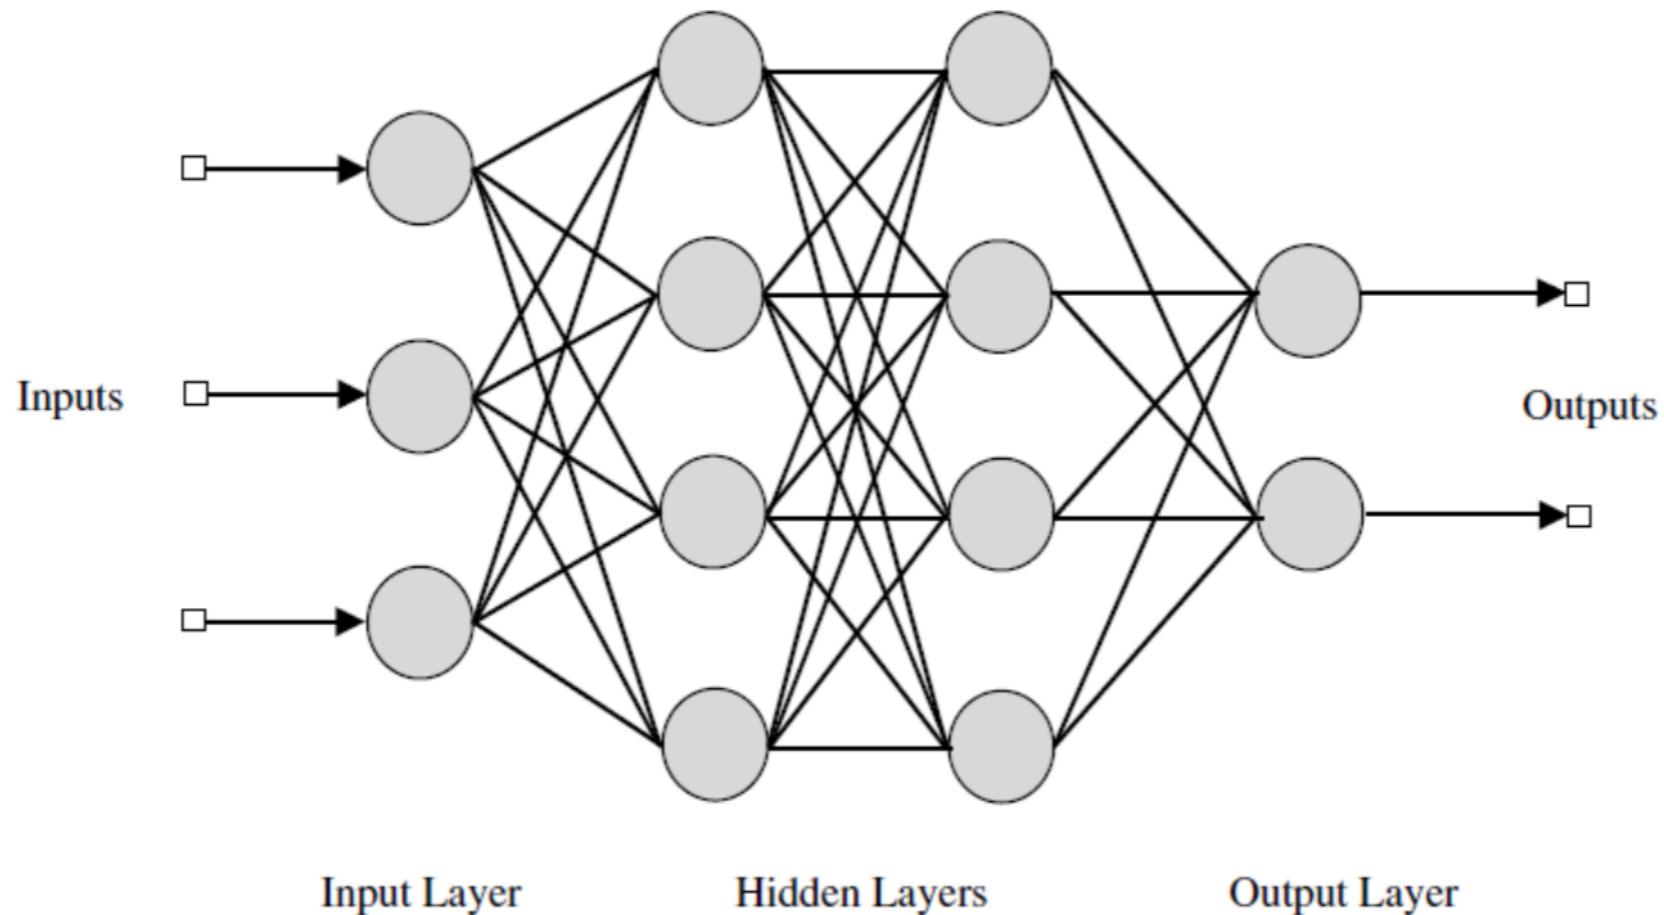
FEED FORWARD NN

- Source



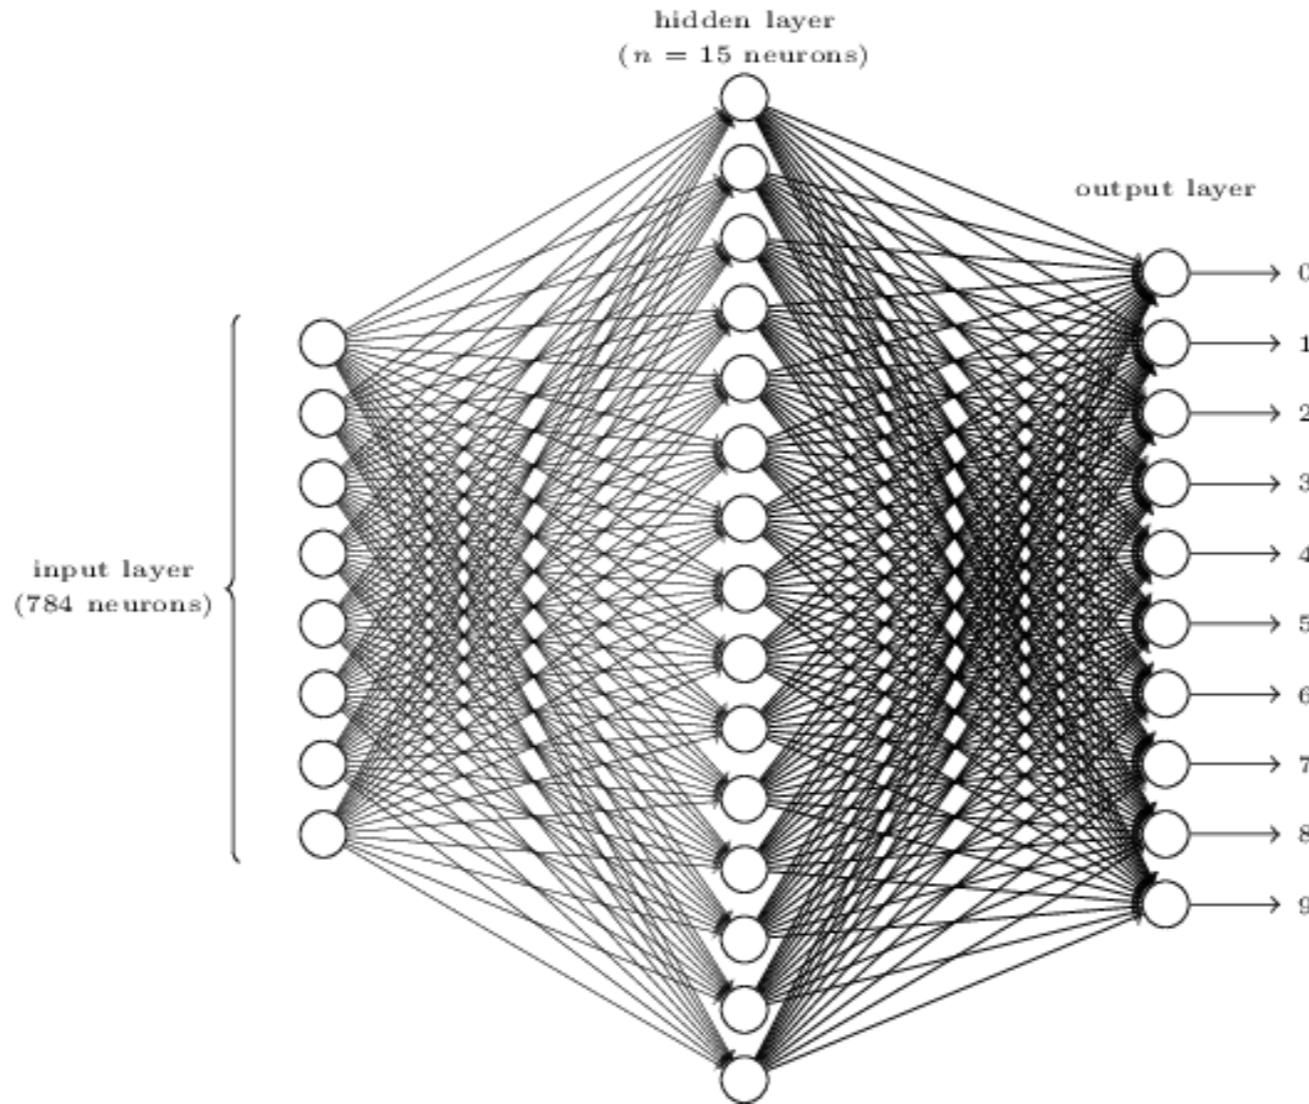
FEED FORWARD NN

▶ Source



FEED FORWARD NN

► [Source](#)



FEED FORWARD NN

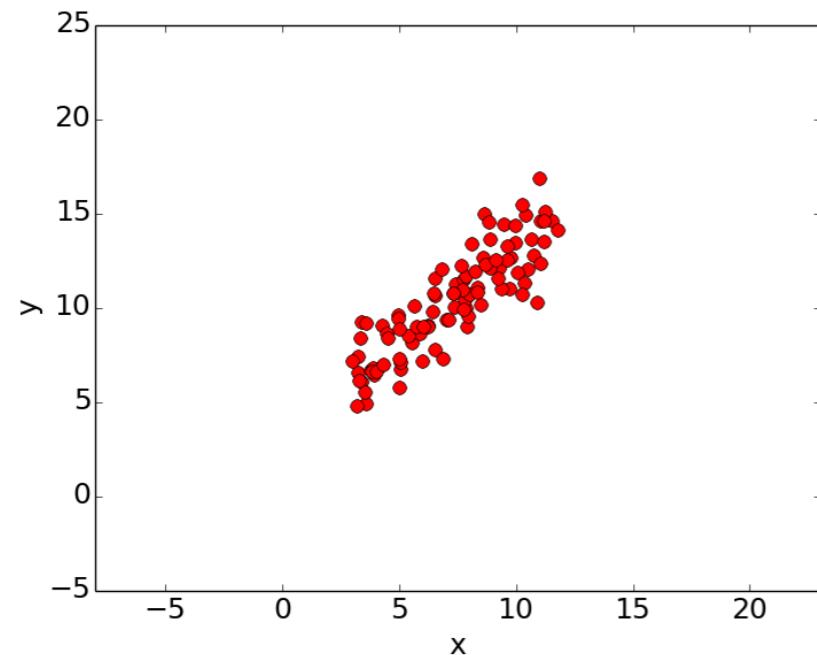
- Typically we use
 - Tanh or logistic layers for input
 - Linear layers for regression output
 - Logistic or Tanh for binary output
 - Softmax for n-class output (yields probabilities)

GUIDED PRACTICE

TRAINING

GRADIENT DESCENT

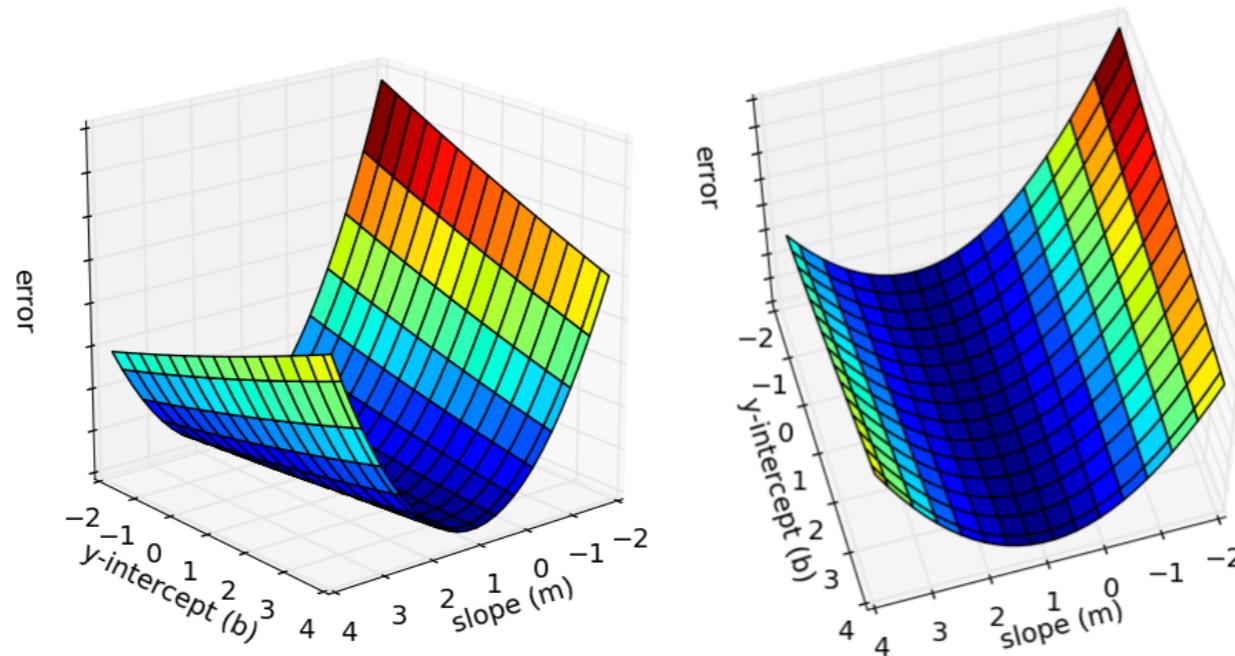
- Goal of Linear Regression is to fit a line to a set of points
- Standard approach is to define an error function (cost function) that measures how good a given line is



$$\text{Error}_{(m,b)} = \frac{1}{N} \sum_{i=1}^N (y_i - (mx_i + b))^2$$

GRADIENT DESCENT

- › Lines that fit the data better will result in lower error values
- › By minimizing the error function we will get the best line for the data
- › When we run gradient descent search, we will start from some location on this surface and move downhill to find the point with the lowest error

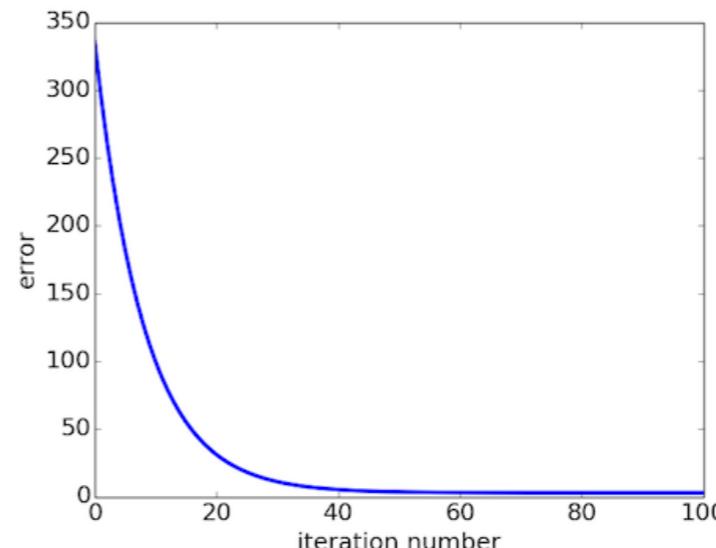


GRADIENT DESCENT

- To run gradient descent on the error function, we need to compute its gradient by differentiating the error function
- Each iteration of the algorithm updates m and b to a line that yields slightly lower error than the previous iteration
- The learning rate controls how large a step to take each iteration
- Iteration stops when error changes are smaller than threshold

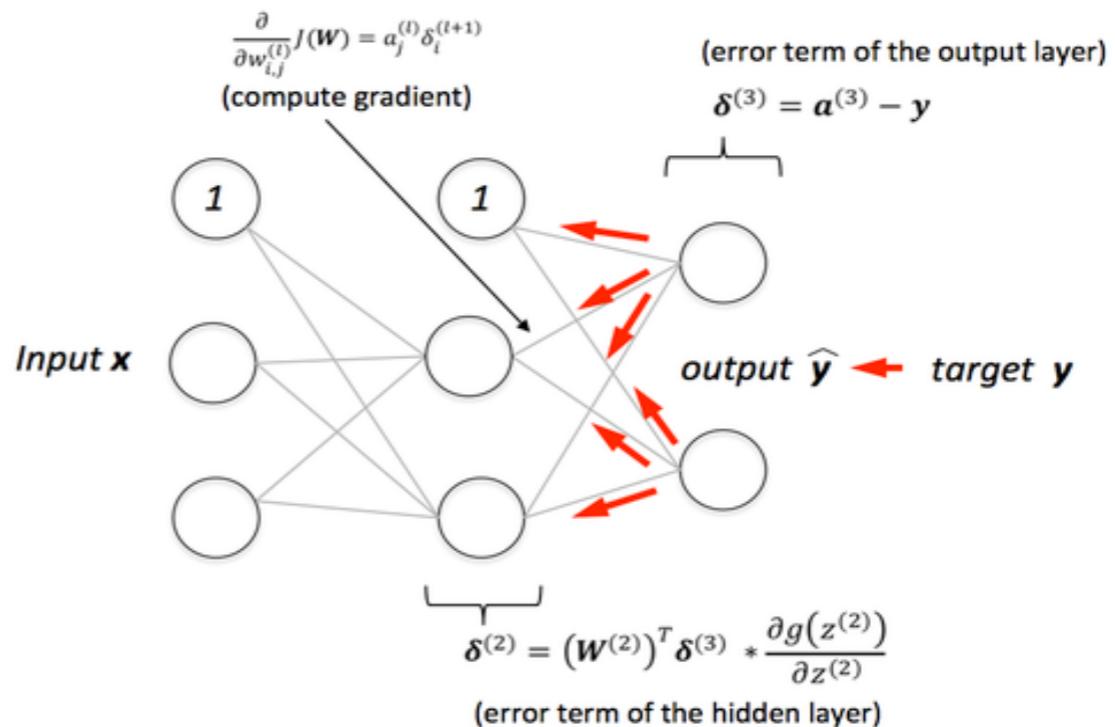
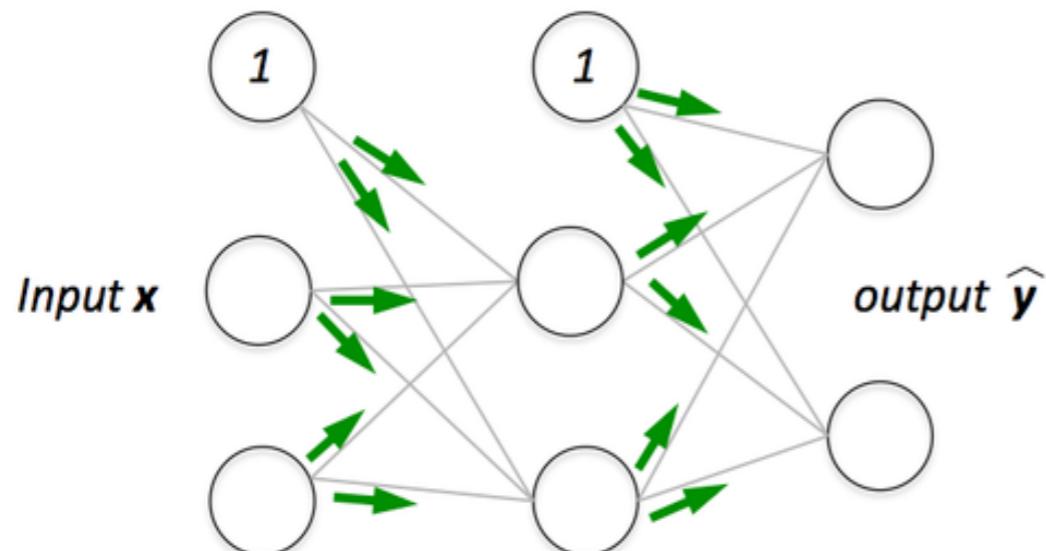
$$\frac{\partial}{\partial m} = \frac{2}{N} \sum_{i=1}^N -x_i(y_i - (mx_i + b))$$

$$\frac{\partial}{\partial b} = \frac{2}{N} \sum_{i=1}^N -(y_i - (mx_i + b))$$



TRAINING

- Feed forward neural networks can be trained with backpropagation
- Source



<http://neuralnetworksanddeeplearning.com/chap2.html>

<https://mattmazur.com/2015/03/17/a-step-by-step-backpropagation-example/>

<https://en.wikipedia.org/wiki/Backpropagation>

TRAINING

- Key Parameters
 - Learning Rate (gradient descent for training)
 - Epochs: number of backpropagation passes (over entire dataset)
 - Batch size: how many training points used at a time to update weights
- Model others behaves as usual with
 - `model.predict`
 - `model.predict_classes`

TRAINING

- Tips
 - If the error jumps around per epoch, decrease the learning rate
 - Taking too long to train: use higher learning rate or batch_size
 - High error after convergence?
 - More hidden layers / neurons
 - Normalize data or use PCA

UNIVERSAL APPROXIMATION

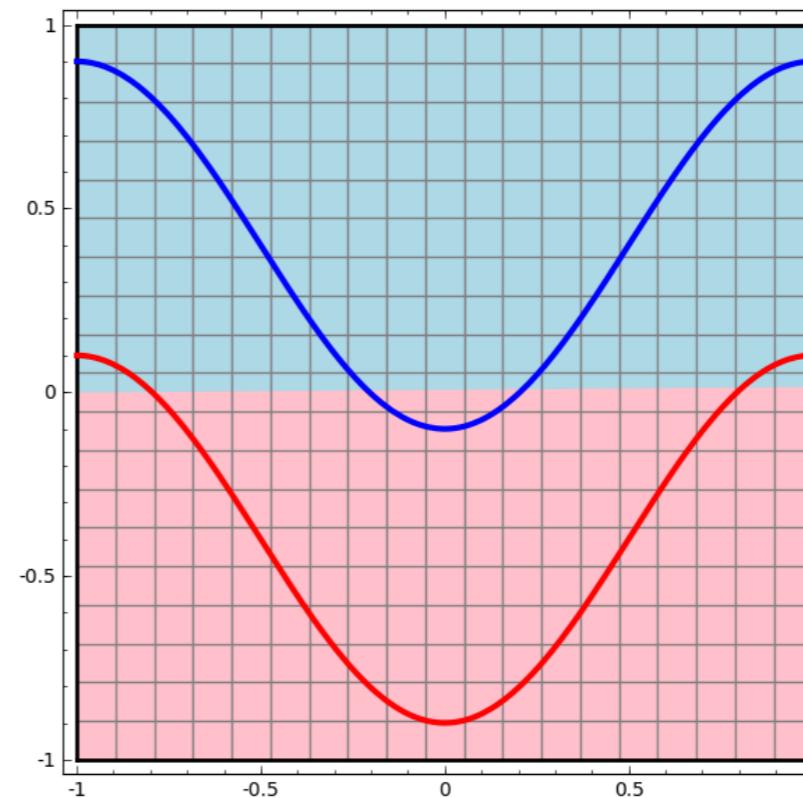
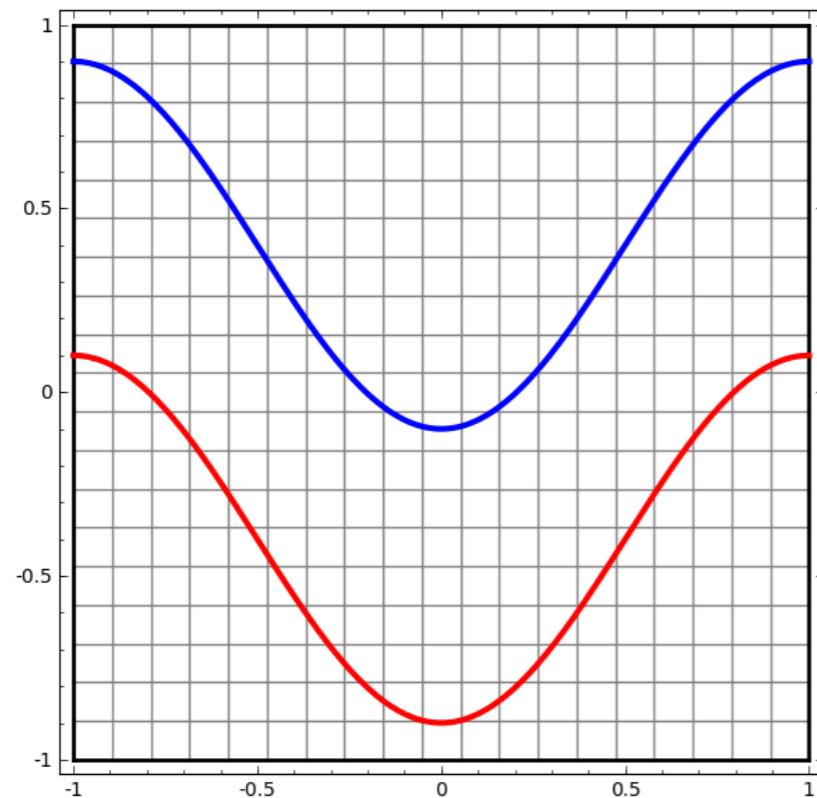
- One major reason that neural networks are useful is the [Universal Approximation Theorem](#)
- The result basically says that many real vector-valued functions can be approximated arbitrarily well with *some* feed-forward neural network
- This is why neural networks are useful for regression -- given enough data and the right network structure they can fit many common data sets

CLASSIFICATION

CLASSIFICATION WITH NEURAL NETWORKS

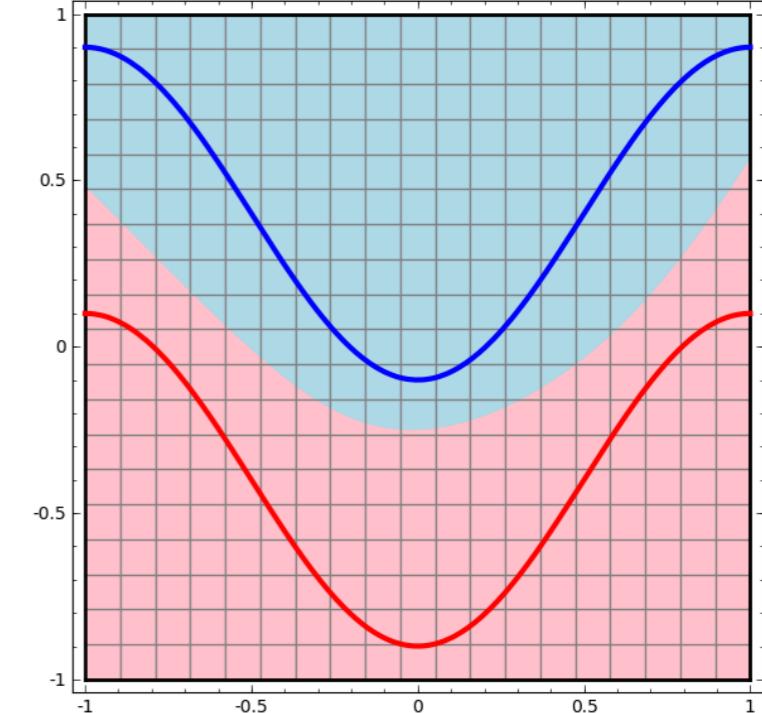
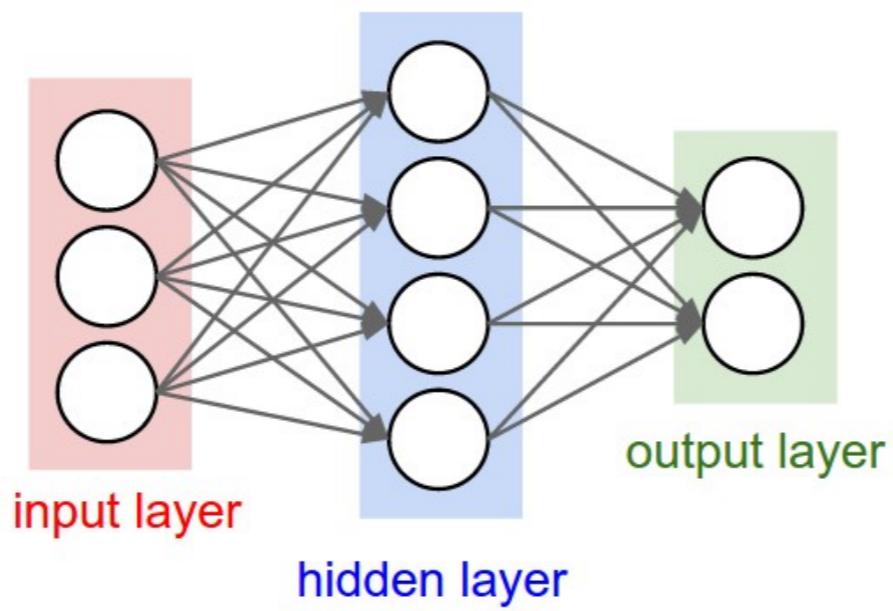
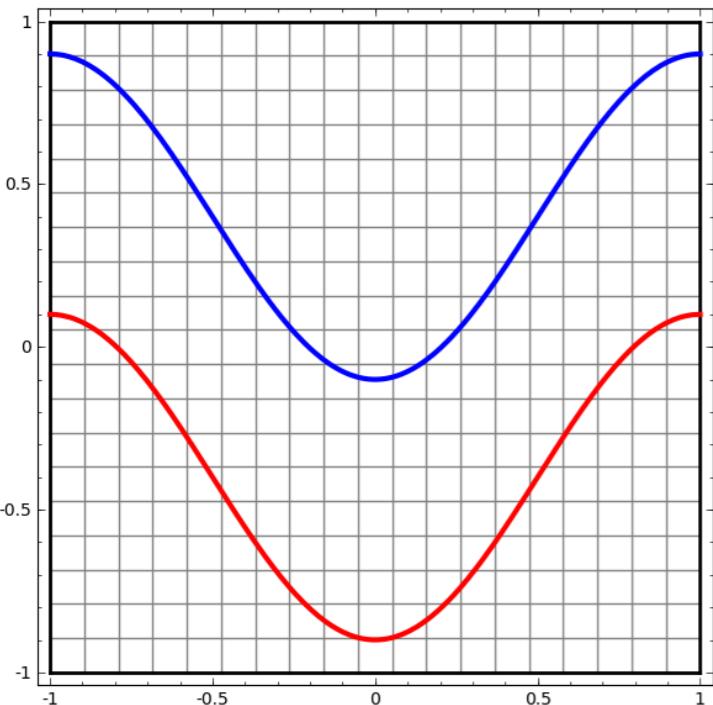
CLASSIFICATION

- Neural Networks are also extremely useful for classification ([source](#))
- No hidden layers:



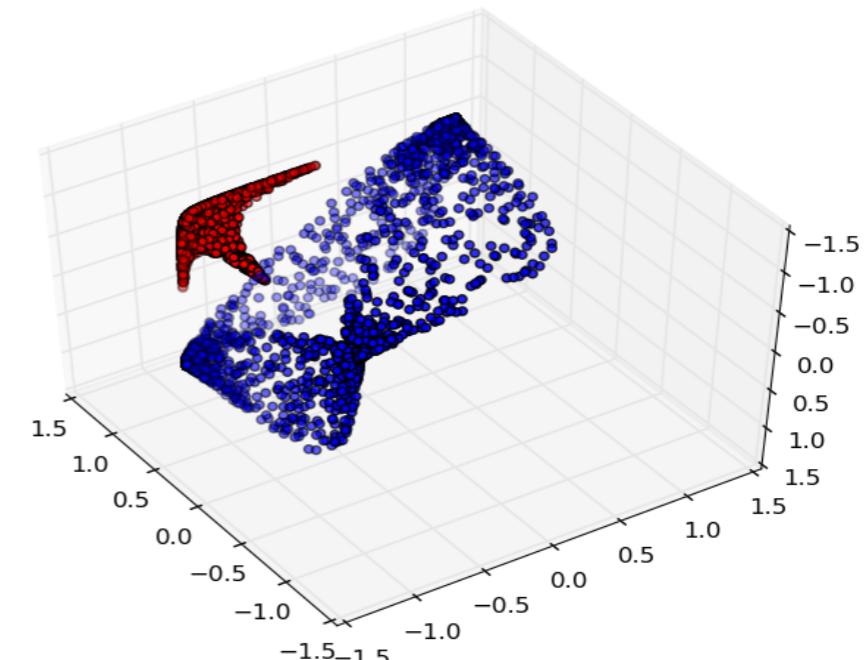
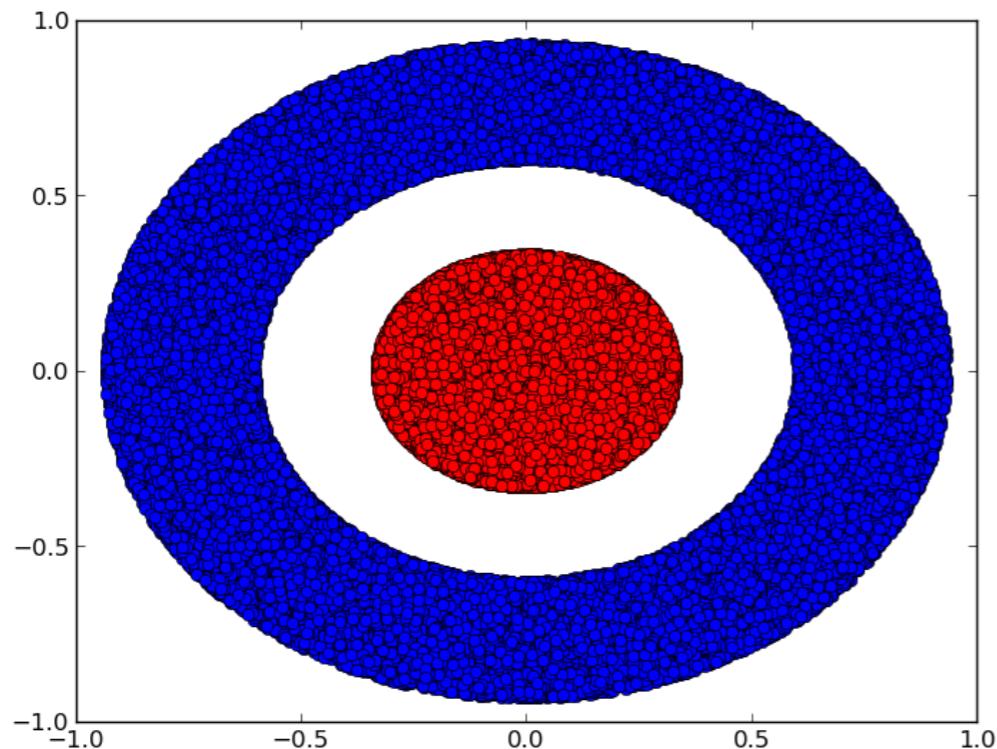
CLASSIFICATION

- Neural Networks are also extremely useful for classification ([source](#))
- One hidden layer:



CLASSIFICATION

- Neural Networks are also extremely useful for classification ([source](#))



CLASSIFICATION

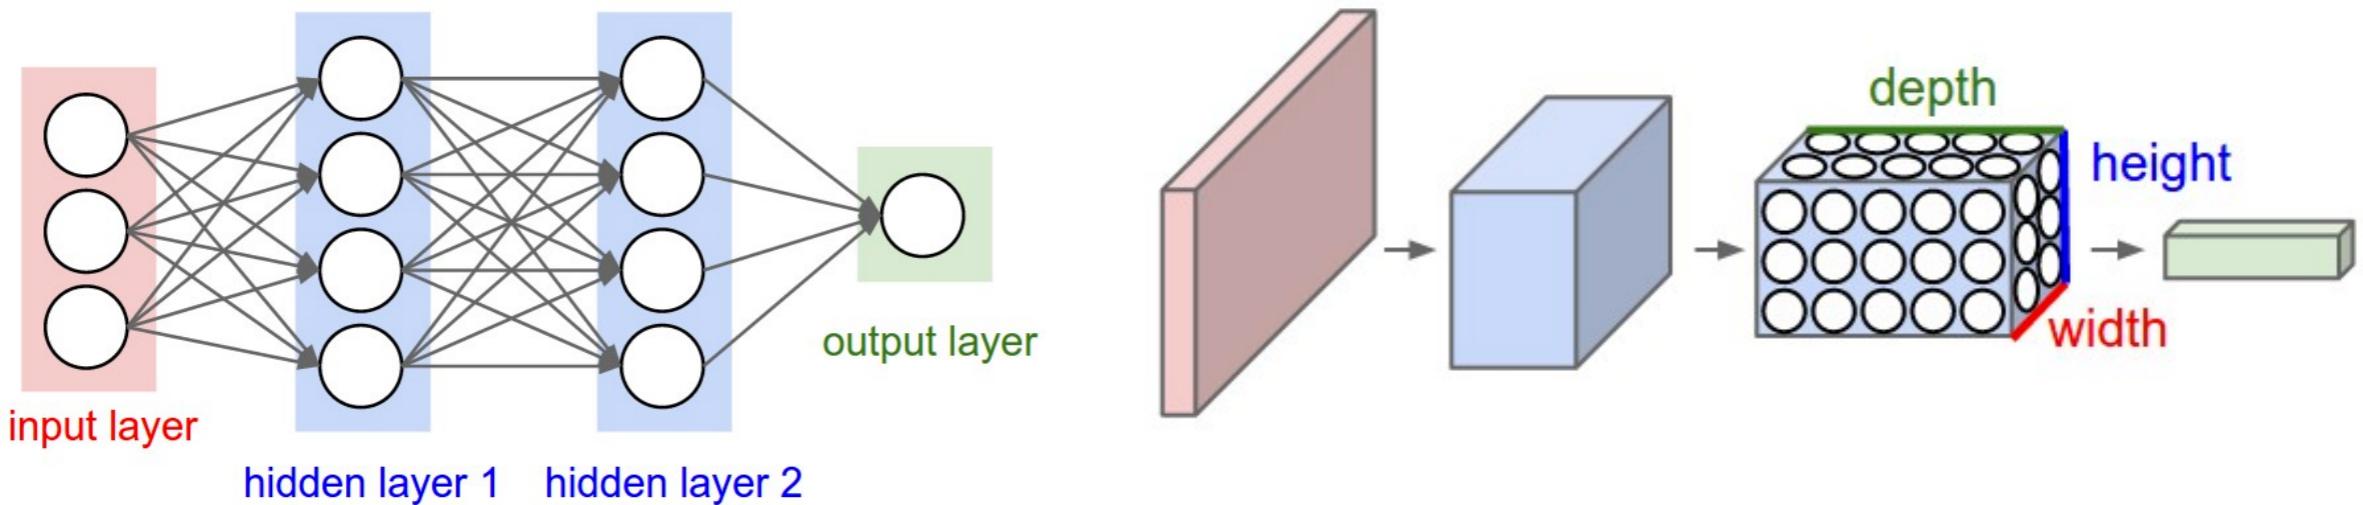
- The neural network transforms the data topologically (no tears or breaks) and then separates the data with a hyperplane
- NNs are capable of handling difficult data sets, including:
 - Image processing: recognizing hand-written characters
 - Image compression
 - Financial forecasting
 - Many others

CONVOLUTIONAL NN

CONVOLUTIONAL NEURAL NETWORKS

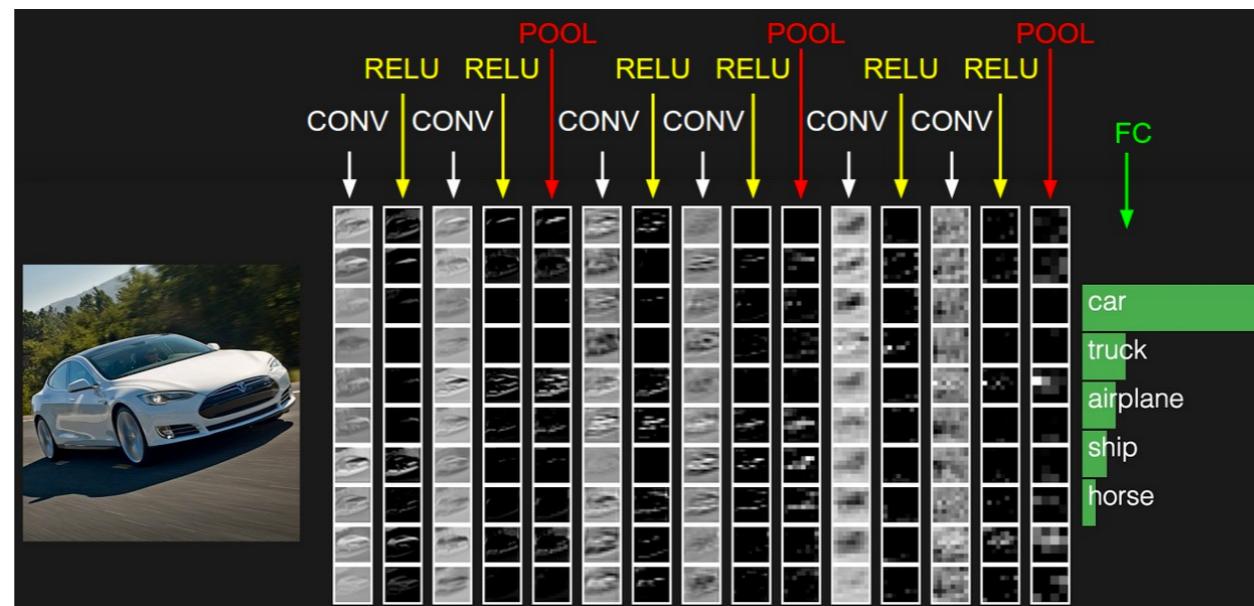
CONVOLUTIONAL NEURAL NETWORKS

- Layers of Convolutional Neural Networks have neurons arranged in 3 dimensions: width, height, and depth
- Every layer of a ConvNet transforms the 3D input volume to a 3D output volume of neuron activations



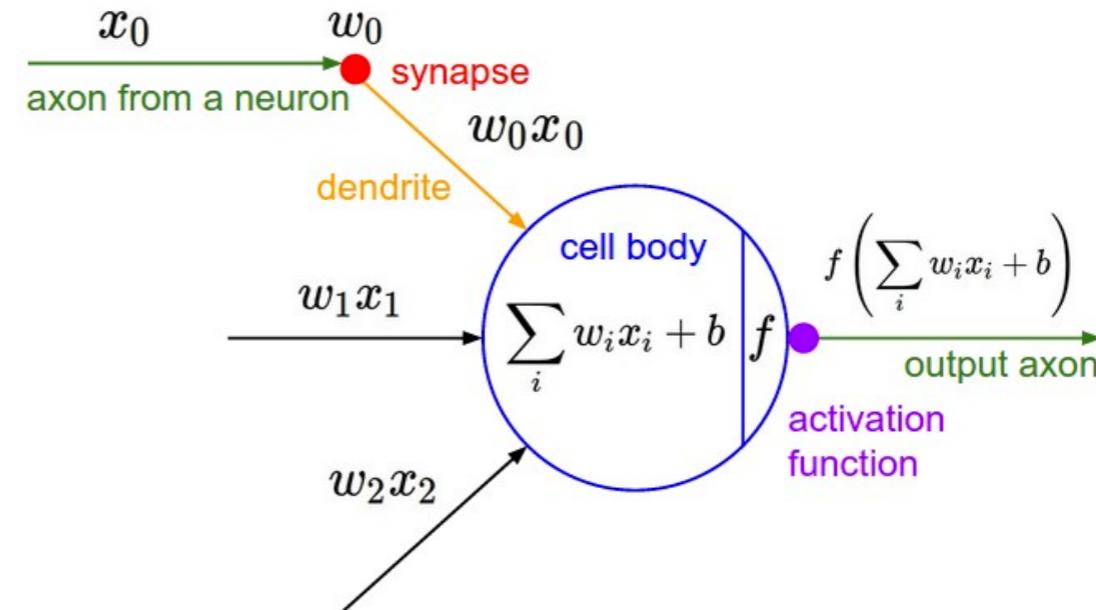
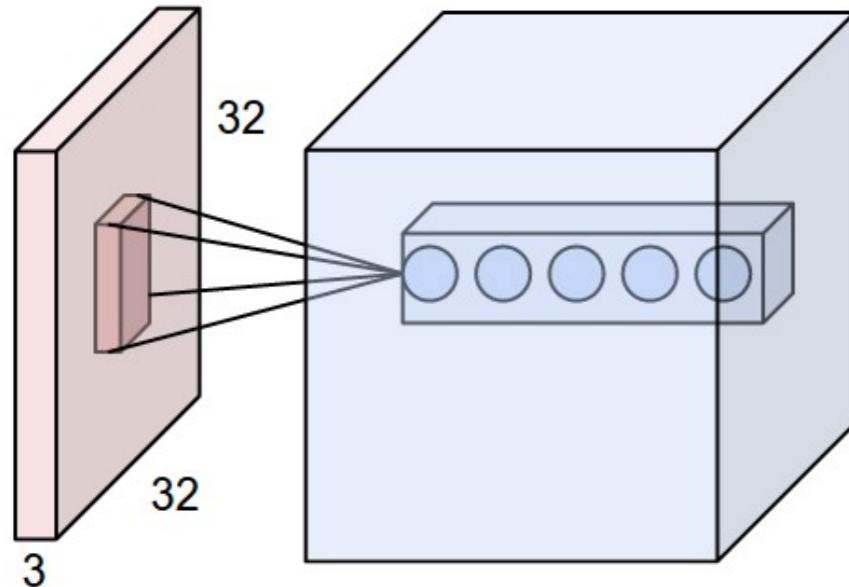
CONVOLUTIONAL NEURAL NETWORKS

- ConvNet architecture represents a list of layers that transform the image volume into an output volume (contains the class scores)
- There are a few distinct types of layers i.e. CONV/RELU/POOL/FC
- Each layer accepts an input 3D volume and transforms it to an output 3D volume through a differentiable function



CONVOLUTIONAL NEURAL NETWORKS

- Each neuron in the convolutional layer is connected only to a local region in the input volume spatially, but to the full depth
- There can be multiple neurons along the depth, all looking at the same region in the input



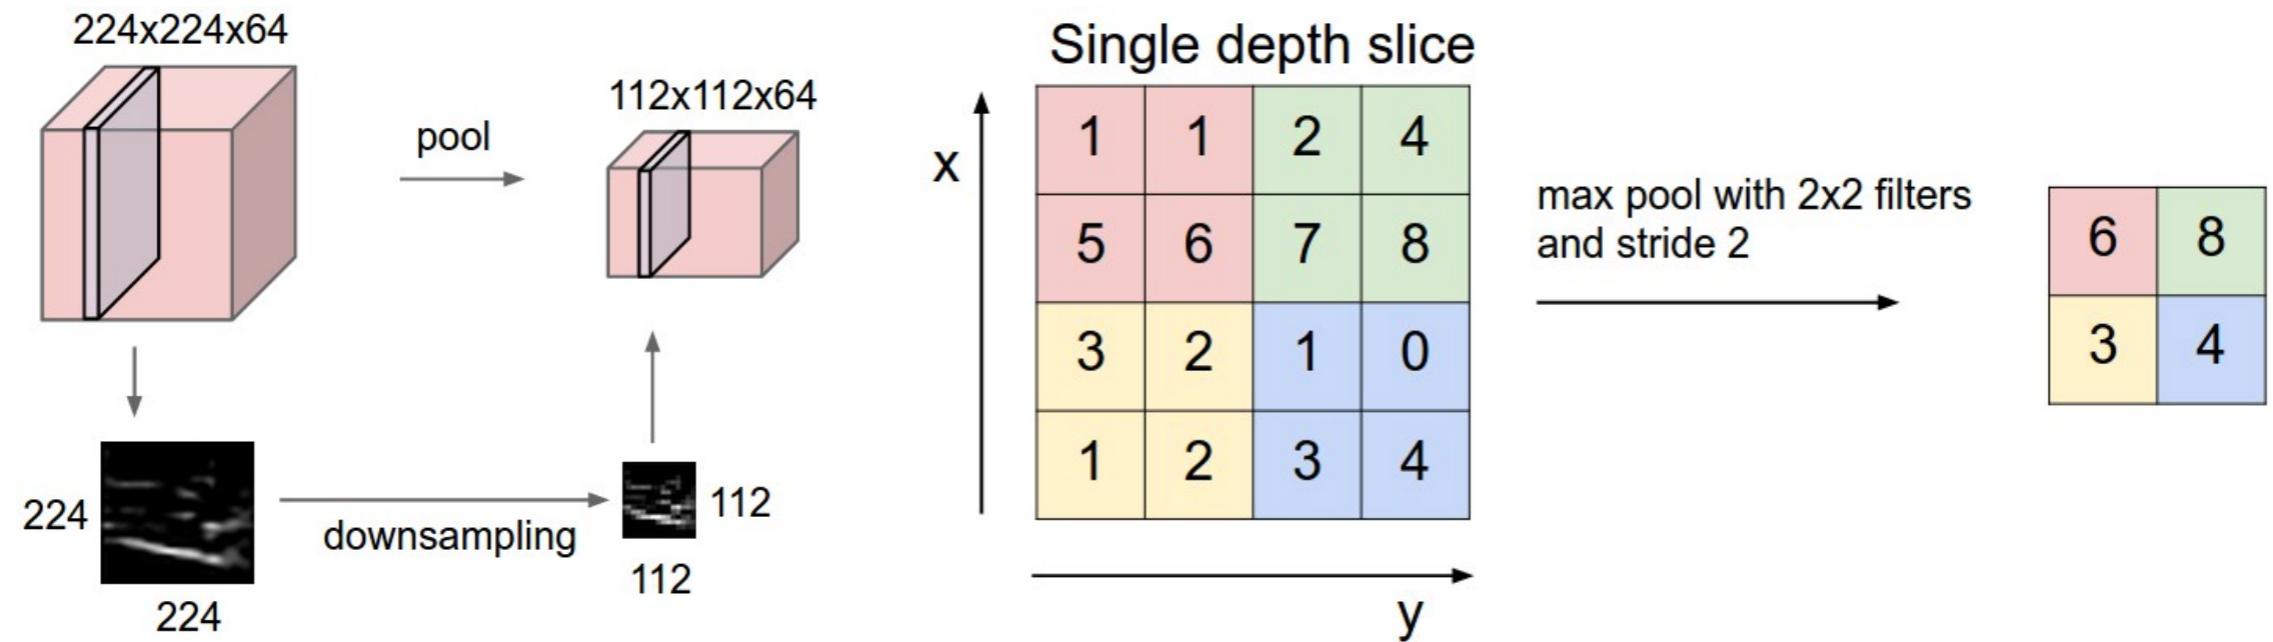
CONVOLUTIONAL NEURAL NETWORKS

- Example filters learned by a ConvNet

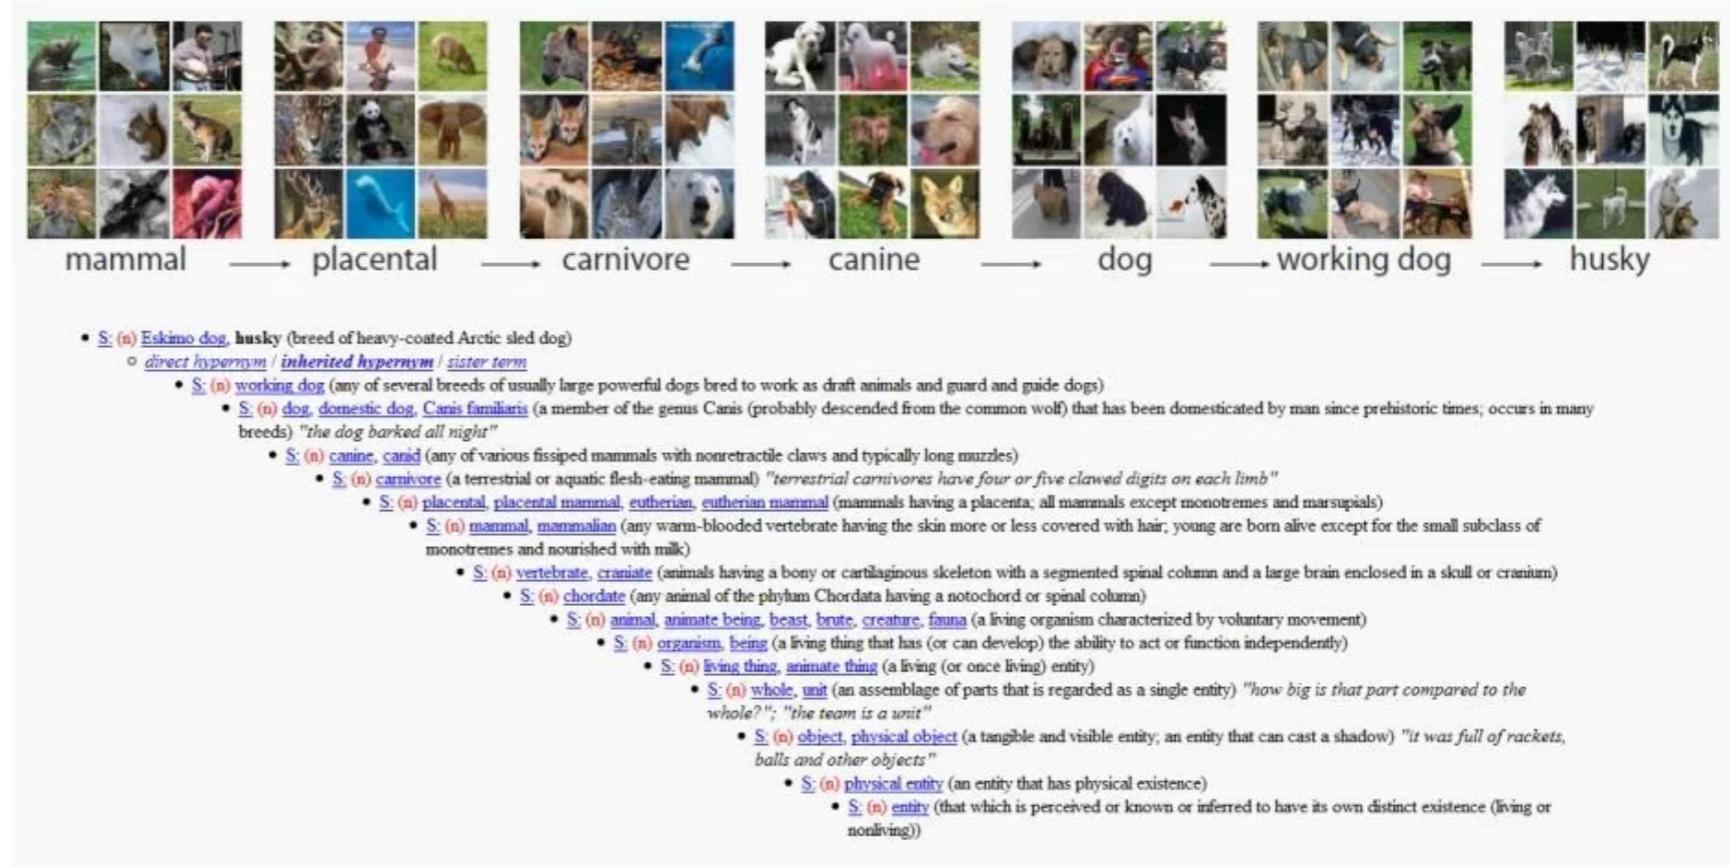
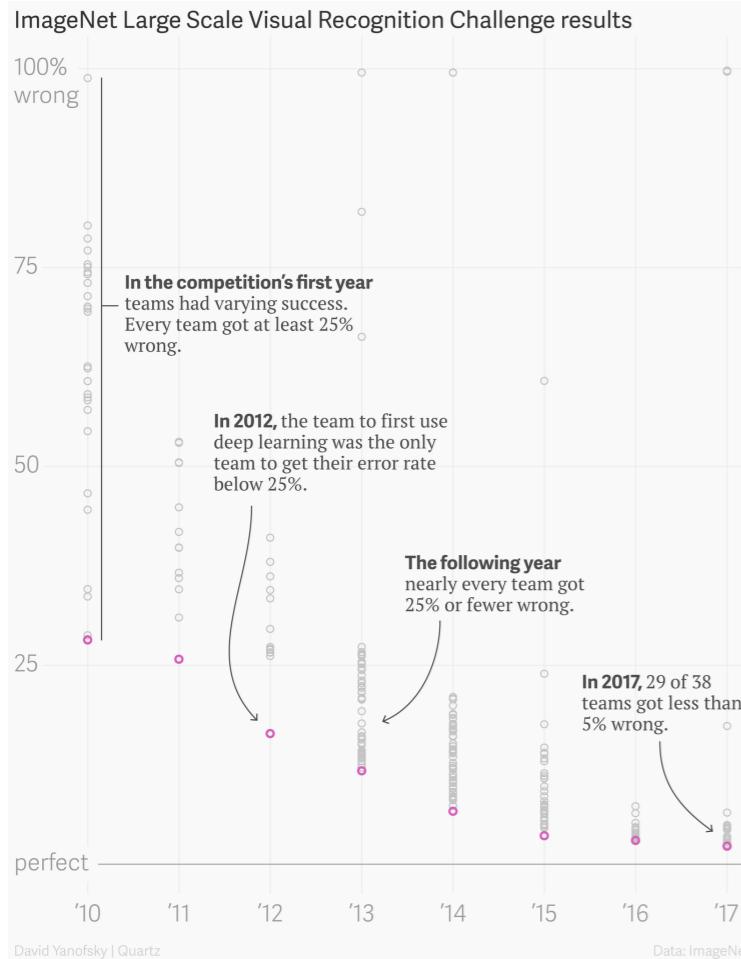


CONVOLUTIONAL NEURAL NETWORKS

- › Pooling layers reduce the spatial size of the representation by down-sampling
- › Max pooling is most commonly used and takes the highest value in each filter region

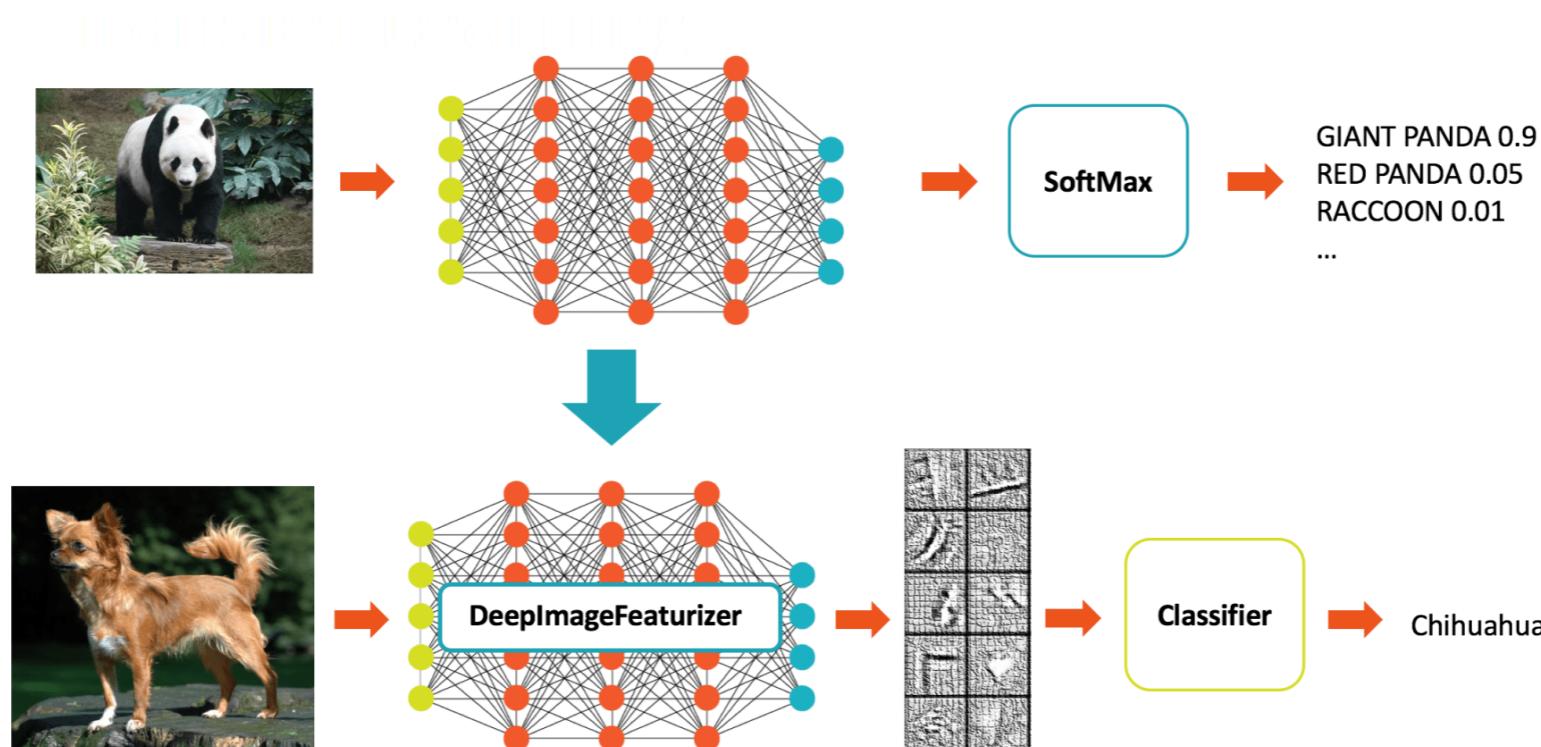


IMAGENET



TRANSFER LEARNING

- In practice, we typically do not need to create our own architecture for a problem but can instead download pre-trained models and fine tune it on our own data

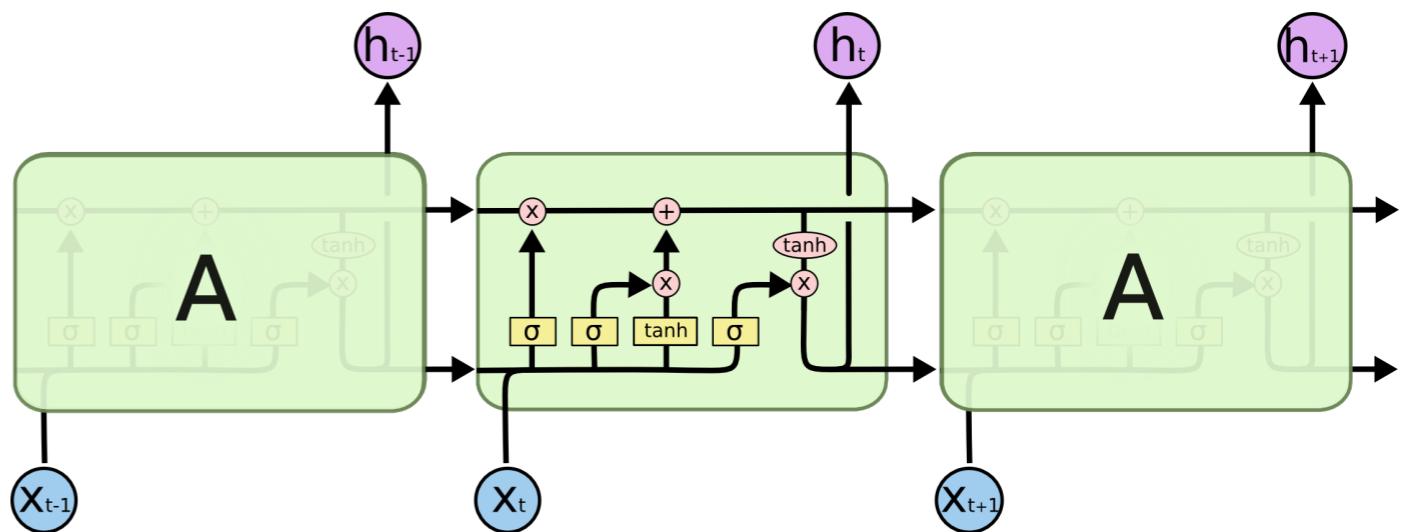
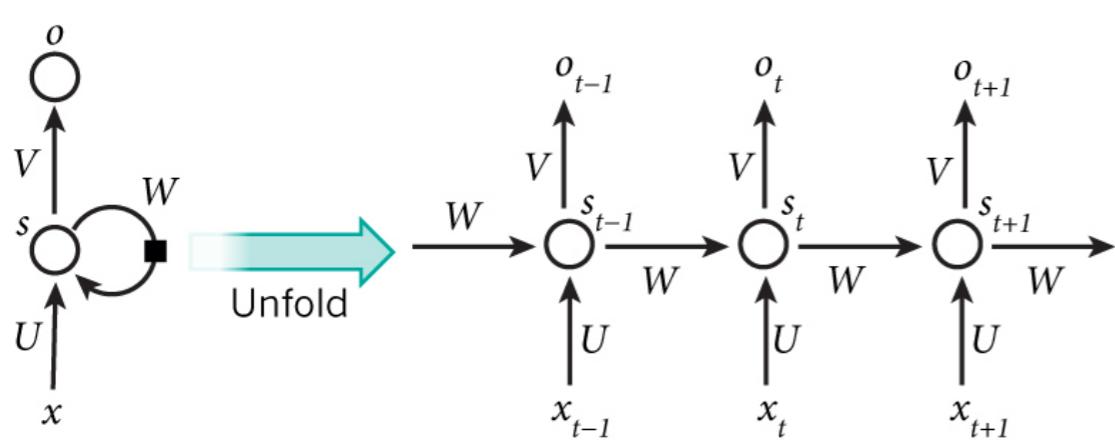


RECURRENT NN

RECURRENT NEURAL NETWORKS

RECURRENT NEURAL NETWORKS

- Recurrent Neural Networks contain loops



<http://www.wildml.com/2015/09/recurrent-neural-networks-tutorial-part-1-introduction-to-rnns/>

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

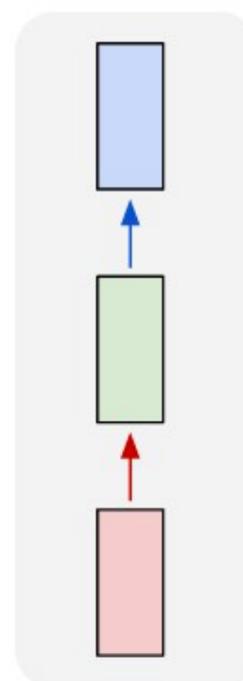
RECURRENT NEURAL NETWORKS

- Recurrent Neural Networks contain loops
- This implements feedback and gives neural networks “memory” or context
- Particularly good for predicting sequences, translating text, recognizing objects in images, speech translation
- Commonly referred to as **deep learning**, involving both feature extraction and modeling

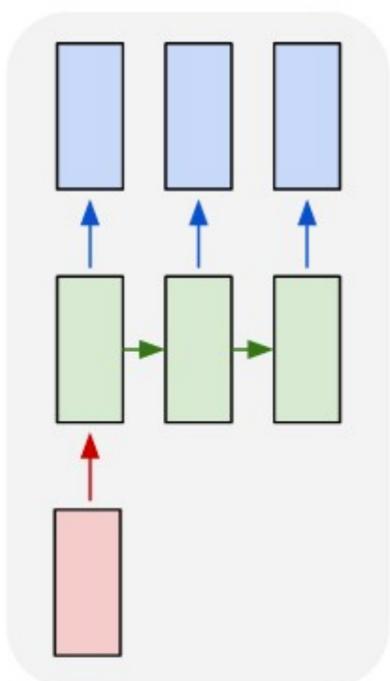
RECURRENT NEURAL NETWORKS

- Recurrent Neural Networks allow us to operate on sequences

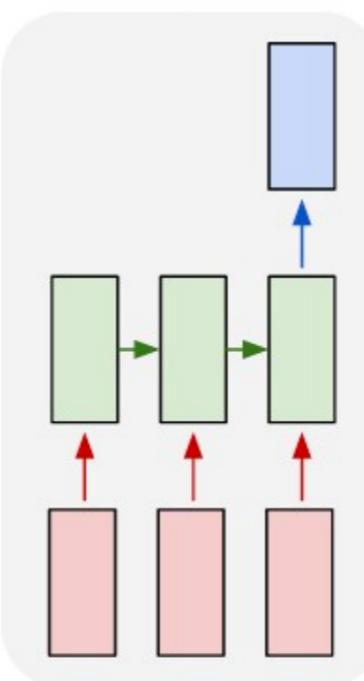
one to one



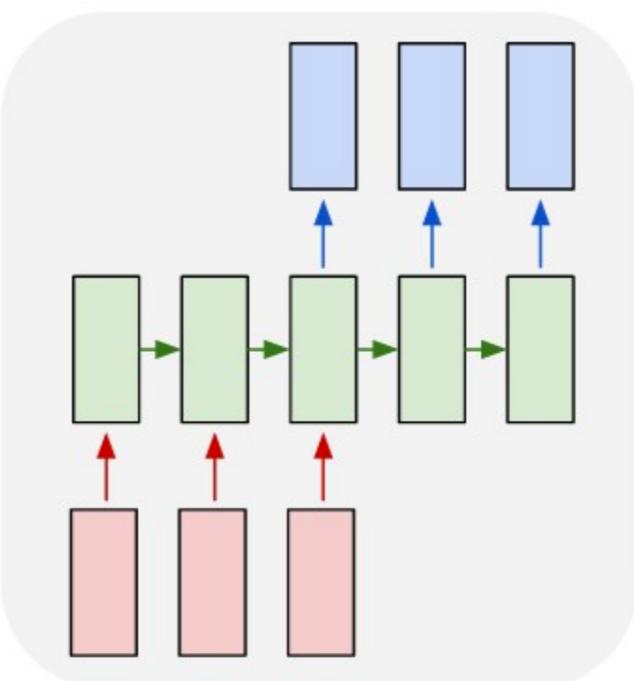
one to many



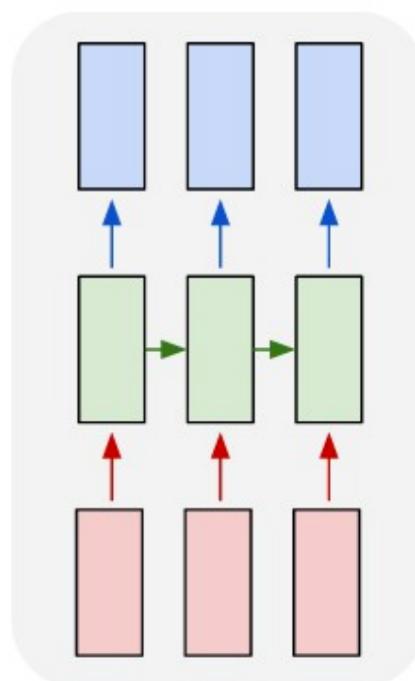
many to one



many to many

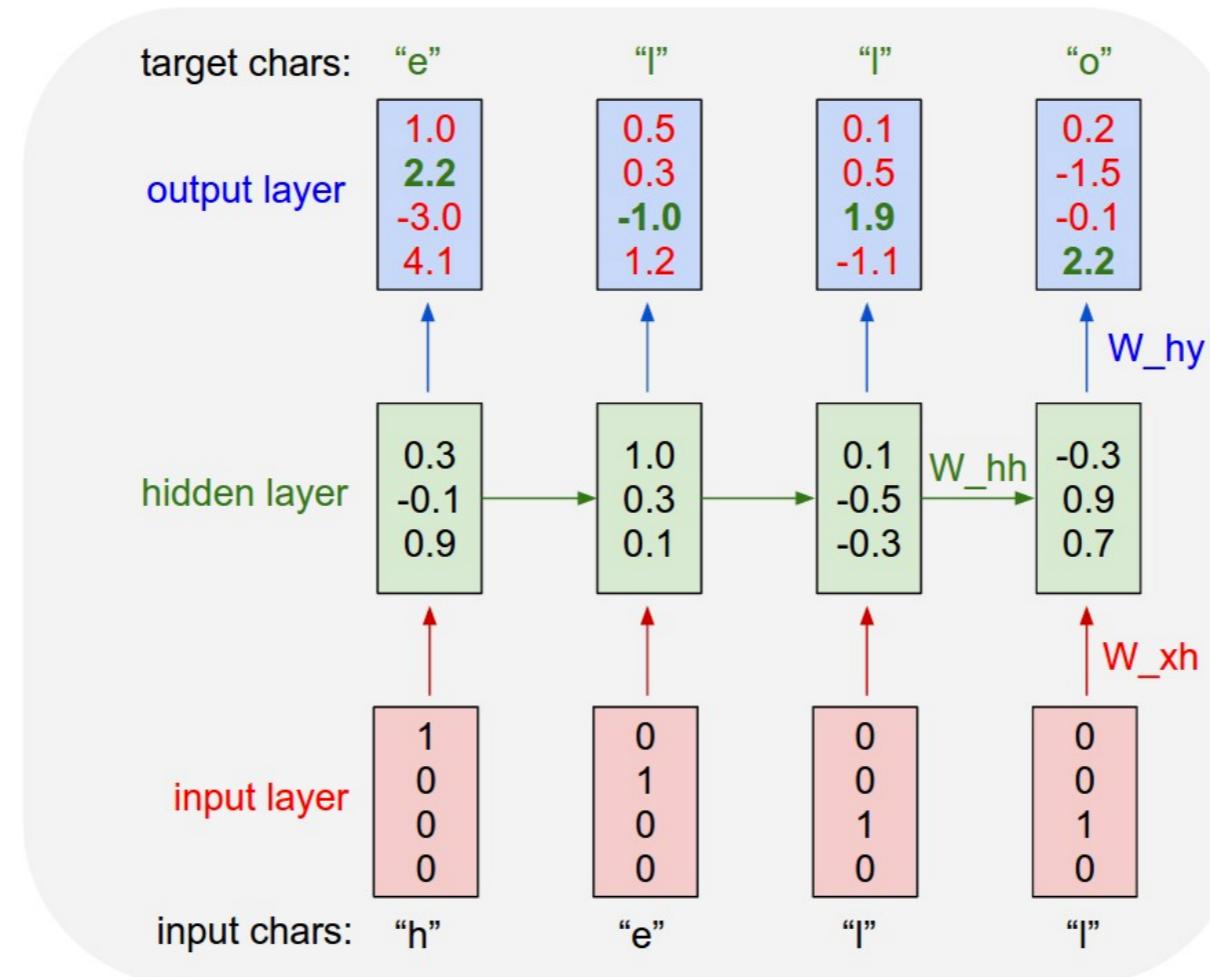


many to many



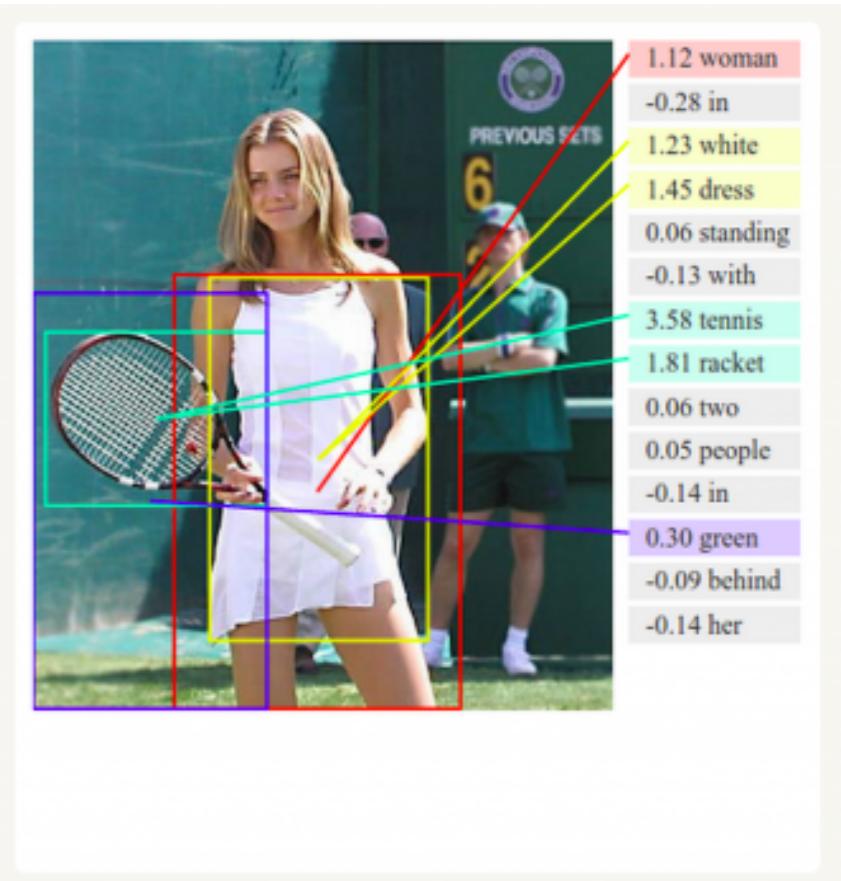
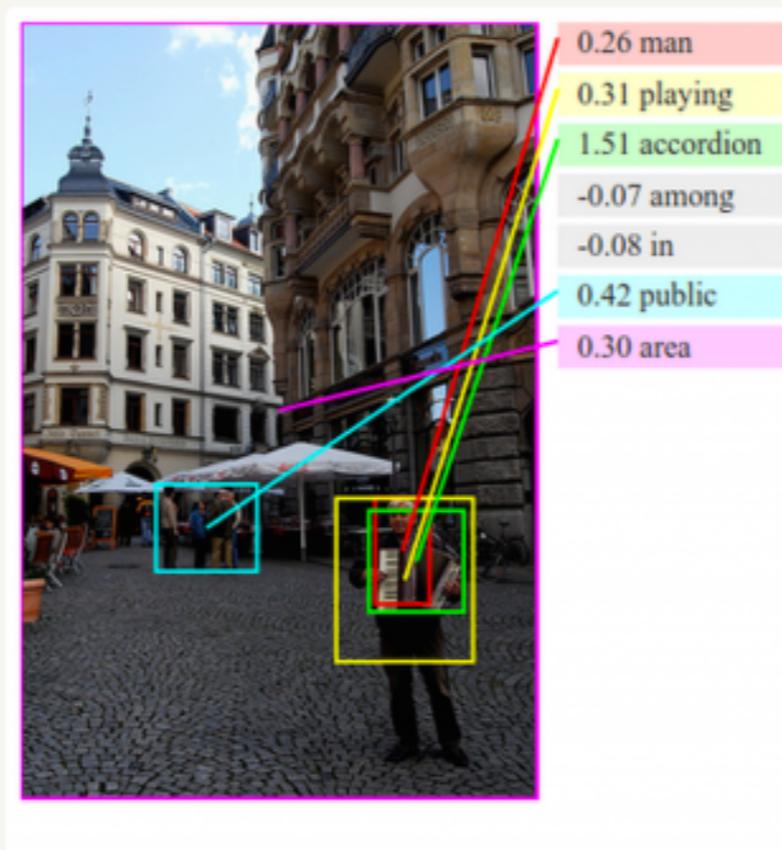
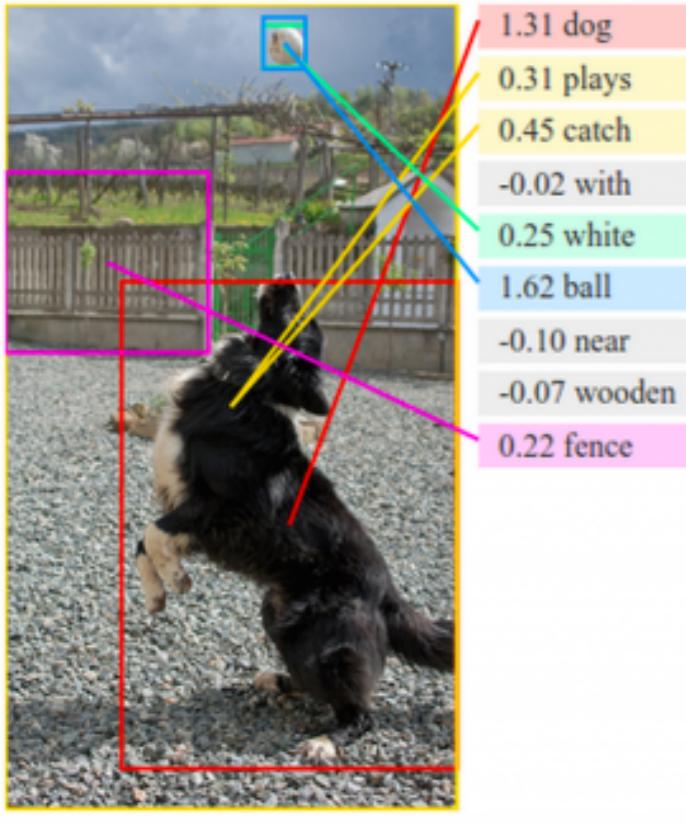
RECURRENT NEURAL NETWORKS

- Training of RNN model



RECURRENT NEURAL NETWORKS

- Generate descriptions for unlabeled images



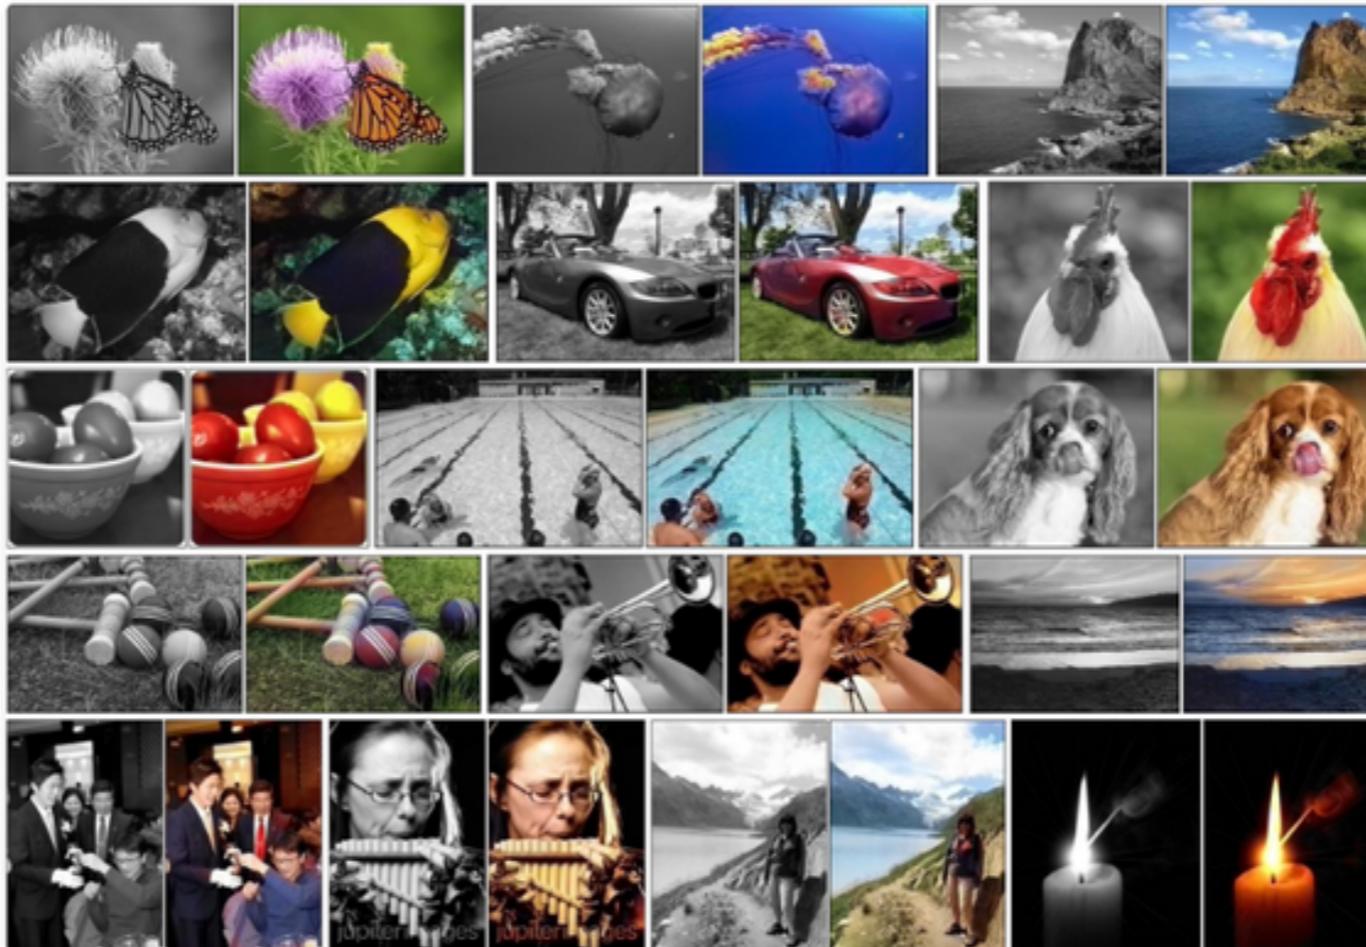
RECURRENT NEURAL NETWORKS

- RNN font analysis

A	B	C	D	E	F	G	H
I	J	K	L	M	N	O	P
O	R	S	T	U	V	W	X
Y	Z	a	b	c	d	e	f
g	h	i	j	k	l	m	n
o	p	q	r	s	t	u	v
w	x	y	z	0	1	2	3
4	5	6	7	8	9		

RECURRENT NEURAL NETWORKS

‣ Automatic Colorization with CNN



<http://tinyclouds.org/colorize/>

GUIDED PRACTICE

NEURAL NETWORKS IN PYTHON

ACTIVITY: KNOWLEDGE CHECK

ANSWER THE FOLLOWING QUESTIONS

EXERCISE

1. Let's practice using [neural networks for classification](#). For each of the four datasets, experiment with the number of layers and neurons to find the best model
2. Also take a look at this [visualization](#)

DELIVERABLE

Answers to the above questions

NN IN PYTHON

- There are many NN libraries for python and other languages
- Python
 - Theano
 - Keras
 - Lasagne
 - TensorFlow
 - Scikit Learn support for NN coming in 0.18
- Lua
 - Torch
- Some of these libraries utilize GPUs for (much) faster training

NN IN PYTHON

- Let's look at some examples in Keras
 - Regression
 - Classification

NN IN PYTHON

- Network design is a hard problem
 - Experience helps
 - Evolutionary algorithms are useful for design
 - Nice (free) book available

CONCLUSION

TOPIC REVIEW

CONCLUSION: Neural Networks

Pros:

- Flexible
- Good for a variety of tasks
- Good for many types of data

Cons:

- Can require a lot of data
- Training may be slow
- Many parameters to tune
- Many layer types and activations
- Black Box model

CONCLUSION

- Many [more examples](#) for Keras available
- Recommended articles: [Convolutional NN](#),
- Advanced machine learning methods you should explore include Bayesian methods and deep learning

<https://github.com/fchollet/keras/tree/master/examples>

<http://cs231n.github.io/neural-networks-1/>

LESSON

Q & A

LESSON

EXIT TICKET

DON'T FORGET TO FILL OUT YOUR EXIT TICKET